



Both effects replicate under four embedders — but only one of them is stable

QASPER — granularity (B→C) (n=290)



HotpotQA — quantifier (A→B) (n=150)



Four models from three training lineages (Sentence-Transformers, BAAI BGE, Microsoft E5). Corpus, questions, chunking and $k=5$ held constant; only the embedder changes. E5 and BGE are asymmetric and are called with their documented query/passage prefixes. p-values are exact McNemar. The granularity gap varies little (14.5-18.3 pp); the quantifier gap nearly halves under the two instruction-trained retrievers.